

## DA Assignment 4

Problem Statement : Abalone Age Prediction

Building a Regression Model

Student Name : PRIYADHARSHINI K

Roll Number : 510119205012

```
# Load the Dataset
```

```
from google.colab import drive  
drive.mount('/content/drive')
```

```
Mounted at /content/drive
```

```
import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
data=pd.read_csv("/content/drive/MyDrive/abalone.csv")
```

```
data.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15

```
data.describe()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
<b>count</b>	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
<b>mean</b>	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	9.933684
<b>std</b>	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
<b>min</b>	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
<b>25%</b>	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	8.000000
<b>50%</b>	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	9.000000
<b>75%</b>	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	11.000000
<b>max</b>	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	29.000000

```
data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sex              4177 non-null   object
1   Length           4177 non-null   float64
2   Diameter         4177 non-null   float64
3   Height           4177 non-null   float64
4   Whole weight     4177 non-null   float64
5   Shucked weight   4177 non-null   float64
6   Viscera weight   4177 non-null   float64
7   Shell weight     4177 non-null   float64
8   Rings            4177 non-null   int64
```

```
dtypes: float64(7), int64(1), object(1)  
memory usage: 293.8+ KB
```

### #Univariate Analysis


```
plt.boxplot(data['Rings'])
```

```
{'whiskers': [<matplotlib.lines.Line2D at 0x7f453e0bac90>,  
<matplotlib.lines.Line2D at 0x7f453e0c4210>],  
'caps': [<matplotlib.lines.Line2D at 0x7f453e0c4750>,  
<matplotlib.lines.Line2D at 0x7f453e0c4c90>],  
'boxes': [<matplotlib.lines.Line2D at 0x7f453e0ba6d0>],  
'medians': [<matplotlib.lines.Line2D at 0x7f453e0cc250>],  
'fliers': [<matplotlib.lines.Line2D at 0x7f453e120c10>],  
'means': []}
```



```
df=pd.DataFrame(data)  
X=list(df.iloc[:,0])  
Y=list(df.iloc[:,1])  
plt.bar(X,Y,color='g')
```


<BarContainer object of 4177 artists>



#Bivariate Analysis


```
plt.bar(data['Diameter'],data['Viscera weight'],color='r')
```

<BarContainer object of 4177 artists>



```
plt.scatter(data['Height'],data['Shell weight'],color='b')
```

```
<matplotlib.collections.PathCollection at 0x7f4537dc8e50>
```




```
# Multivariate Analysis
```

```
sns.catplot(data=data,x="Height",y="Shucked weight",hue="Sex")
```

```
<seaborn.axisgrid.FacetGrid at 0x7f4537bbe2d0>
```

```
sns.catplot(data=data,x="Whole weight",y="Shell weight",hue="Sex",kind="box")
```

```
<seaborn.axisgrid.FacetGrid at 0x7f4538a19ed0>
```



```
# Perform Descriptive Statistics on the Dataset
```

```
data.mean()
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reduction
  """Entry point for launching an IPython kernel.
Length          0.523992
Diameter        0.407881
Height          0.139516
Whole weight    0.828742
Shucked weight  0.359367
Viscera weight  0.180594
```

```

Shell weight    0.238831
Rings           9.933684
dtype: float64

```

```
data.median()
```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reduct
"""Entry point for launching an IPython kernel.
Length          0.5450
Diameter         0.4250
Height          0.1400
Whole weight     0.7995
Shucked weight   0.3360
Viscera weight   0.1710
Shell weight     0.2340
Rings           9.0000
dtype: float64

```

```
data.describe()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
<b>count</b>	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
<b>mean</b>	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	9.933684
<b>std</b>	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
<b>min</b>	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
<b>25%</b>	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	8.000000
<b>50%</b>	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	9.000000
<b>75%</b>	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	11.000000
<b>max</b>	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	29.000000

```
data.shape
```

```
(4177, 9)
```

```
# Handling the Missing Values
```

```
data.isna().sum()
```

```
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Shucked weight 0
Viscera weight 0
Shell weight 0
Rings        0
dtype: int64
```

```
# Find the Outliers and Replace the Outliers
```

```
sns.boxplot(data['Height'])
```



```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg:
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7f4539b0ddd0>
```



```
qnt=data.quantile(q=[0.25,0.75])
qnt
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
<b>0.25</b>	0.450	0.35	0.115	0.4415	0.186	0.0935	0.130	8.0
<b>0.75</b>	0.615	0.48	0.165	1.1530	0.502	0.2530	0.329	11.0

Height

```
IQR=qnt.loc[0.75] - qnt.loc[0.25]
IQR
```

```
Length      0.1650
Diameter    0.1300
Height      0.0500
Whole weight 0.7115
Shucked weight 0.3160
Viscera weight 0.1595
Shell weight 0.1990
Rings       3.0000
dtype: float64
```

```
upper_extreme=qnt.loc[0.75]+1.5*IQR
upper_extreme
```

```
Length      0.86250
Diameter    0.67500
Height      0.24000
Whole weight 2.22025
Shucked weight 0.97600
```

```
Viscera weight    0.49225  
Shell weight      0.62750  
Rings             15.50000  
dtype: float64
```

```
lower_extreme=qnt.loc[0.25]-1.5*IQR  
lower_extreme
```

```
Length           0.20250  
Diameter         0.15500  
Height           0.04000  
Whole weight     -0.62575  
Shucked weight   -0.28800  
Viscera weight   -0.14575  
Shell weight     -0.16850  
Rings            3.50000  
dtype: float64
```

```
df=data[(data['Height']<upper_extreme['Height'])&(data['Height']>lower_extreme['Height'])]
```

```
data.shape
```


```
(4177, 9)
```

```
df.shape
```

```
(4148, 9)
```

```
sns.boxplot(df['Height'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg:
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7f453932c750>
```




```
# Check for Categorical Columns and Perform Encoding
from sklearn.preprocessing import LabelEncoder

Height
```

```
le=LabelEncoder()
df['Sex']=le.fit_transform(df['Sex'])
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-vs-returning-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-vs-returning-a-copy)



```
df.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15

# Split the Data into Dependent and Independent Variables

```
2    0    0.530    0.420    0.135    0.6770    0.2565    0.1415    0.210    9
```

```
y=df['Rings']
```

```
x=df.drop(columns=['Rings'],axis=1)
```

```
4    1    0.330    0.255    0.080    0.2050    0.0895    0.0395    0.055    1
```

```
x.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055

# Scale the Independent Variables

```
from sklearn.preprocessing import scale
```

```
names = x.columns
```

```
names
```

```
Index(['Sex', 'Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',  
      'Viscera weight', 'Shell weight'],  
      dtype='object')
```

```
x=scale(x)
```

```
x
```

```
array([[ 1.1486969 , -0.60262947, -0.45591118, ..., -0.61827733,
        -0.73839808, -0.64915159],
       [ 1.1486969 , -1.49739064, -1.48593272, ..., -1.18589191,
        -1.22152247, -1.22814023],
       [-1.27860557,  0.03648565,  0.11060067, ..., -0.472968  ,
        -0.36570212, -0.2149101  ],
       ...,
       [ 1.1486969 ,  0.63299309,  0.67711252, ...,  0.74853858,
        0.97784382,  0.49435099],
       [-1.27860557,  0.84603146,  0.78011468, ...,  0.77351362,
        0.73398103,  0.40750269],
       [ 1.1486969 ,  1.57036193,  1.50112976, ...,  2.65572357,
        1.7968547  ,  1.84773695]])
```

```
x=pd.DataFrame(x,columns=[names])
```

```
x.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight
0	1.148697	-0.602629	-0.455911	-1.194353	-0.653289	-0.618277	-0.738398	-0.649152
1	1.148697	-1.497391	-1.485933	-1.327737	-1.246310	-1.185892	-1.221522	-1.228140
2	-1.278606	0.036486	0.110601	-0.127275	-0.318237	-0.472968	-0.365702	-0.214910
3	1.148697	-0.730452	-0.455911	-0.394044	-0.649178	-0.659146	-0.618767	-0.612965
4	-0.064954	-1.667821	-1.588935	-1.594507	-1.288449	-1.231301	-1.304344	-1.336701

```
# Split the data into training and testing
```

```
from sklearn.model_selection import train_test_split
```

```
df.iloc[:, :-1]
df.iloc[:, -1]
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.05, random_state=0)
```

x\_train

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight
<b>531</b>	-0.064954	-0.048730	-0.043903	0.006110	-0.522763	-0.729530	-0.333494	-0.439268
<b>2973</b>	-1.278606	0.846031	0.831616	-0.794199	0.625252	0.753079	0.094416	0.382172
<b>528</b>	-1.278606	-0.773060	-0.764918	-0.794199	-0.928730	-0.868028	-0.876434	-0.938646
<b>2062</b>	-1.278606	0.803424	0.831616	0.539649	0.460809	0.600959	0.577541	0.237425
<b>991</b>	-1.278606	0.590385	0.419607	0.272879	0.580030	1.025534	0.361285	0.074584
...	...	...	...	...	...	...	...	...
<b>1033</b>	-1.278606	1.272108	1.604132	1.473342	2.067208	1.833818	1.658819	2.267504
<b>3264</b>	1.148697	0.079093	0.110601	0.539649	-0.229849	-0.388961	-0.158649	-0.106350
<b>1653</b>	-1.278606	0.675601	0.986119	1.073187	1.237801	1.670345	0.761588	1.384546
<b>2607</b>	-0.064954	-1.454783	-1.279928	-1.194353	-1.242199	-1.204056	-1.230725	-1.217284
<b>2732</b>	-0.064954	-0.006122	-0.095404	-0.260660	-0.385042	-0.071097	-0.485333	-0.287284

3940 rows × 8 columns

x\_test

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight
<b>825</b>	-0.064954	-0.687845	-0.867920	-1.060968	-0.932842	-0.822619	-0.747600	-1.090630
<b>4077</b>	-1.278606	1.655577	1.861637	1.740111	1.886321	1.883768	2.220164	1.739177
<b>971</b>	-1.278606	0.206916	0.368106	0.406264	0.326172	0.832546	-0.029816	-0.157011
<b>1499</b>	1.148697	1.059070	0.780115	0.006110	0.705418	0.519222	0.572940	-0.178723
<b>3176</b>	-1.278606	-1.199137	-1.125425	-1.060968	-1.149700	-1.131401	-1.143302	-1.155767
...	...	...	...	...	...	...	...	...
<b>1927</b>	1.148697	0.888639	0.831616	0.406264	0.864722	1.223064	0.595945	0.356841
<b>2052</b>	1.148697	0.206916	0.007599	-0.260660	0.079508	0.385265	0.278464	-0.193198
<b>3169</b>	-1.278606	0.718208	0.625611	0.539649	0.380644	0.296717	0.407297	0.581199

y\_train

```

536      11
2997      9
533       9
2079      9
1000     11
      ..
1042     12
3289     15
1667     11
2630      6
2756      8

```

Name: Rings, Length: 3940, dtype: int64

y\_test

```

834      7
4106     11
980      9

```

```
1513      8
3201      6
..
1942      9
2069      9
3194     12
998       8
604       12
Name: Rings, Length: 208, dtype: int64
```

```
# Build the Model
```


```
from sklearn.linear_model import LinearRegression
```

```
lin_reg=LinearRegression()
```

```
# Train the Model
```

```
lin_reg.fit(x,y)
```


```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/validation.py:1692: FutureWarning: Feature names only support names that a
FutureWarning,
LinearRegression()
```



```
# Test the Model
```

```
lin_reg.predict(x)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/validation.py:1692: FutureWarning: Feature names only support names that a
FutureWarning,
array([ 8.50595288,  7.31132061, 10.72922803, ..., 11.51826065,
        9.37712194, 11.10957517])
```






```
predict=lin_reg.predict(x) predict
```

```
sns.scatterplot(df['Diameter'],df['Rings'],color='red')
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args:
FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7f4532dd79d0>
```



```
def viz_linear():
    sns.scatter(x,y,color='red')
    sns.plot(x,prediction,color='blue')

viz_linear()
```

```
from sklearn.preprocessing import PolynomialFeatures
poly_reg=PolynomialFeatures(degree=2)
x_poly=poly_reg.fit_transform(x)
x_poly
```

```

/usr/local/lib/python3.7/dist-packages/sklearn/utils/validation.py:1692: FutureWarning: Feature names only support names that a
FutureWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/utils/validation.py:1692: FutureWarning: Feature names only support names that a
FutureWarning,
array([[ 1.          ,  1.1486969 , -0.60262947, ...,  0.54523172,
         0.47933228,  0.42139778],
       [ 1.          ,  1.1486969 , -1.49739064, ...,  1.49211715,
         1.50020089,  1.50832843],
       [ 1.          , -1.27860557,  0.03648565, ...,  0.13373804,
         0.07859308,  0.04618635],
       ...,
       [ 1.          ,  1.1486969 ,  0.63299309, ...,  0.95617853,
         0.48339806,  0.2443829 ],
       [ 1.          , -1.27860557,  0.84603146, ...,  0.53872815,
         0.29909925,  0.16605844],
       [ 1.          ,  1.1486969 ,  1.57036193, ...,  3.2286868 ,
         3.32011482,  3.41413183]])

```

```

pol_reg=LinearRegression()
pol_reg.fit(x_poly,y)
prediction=pol_reg.predict(x_poly)

```

```
df.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

```
# Performance using Evaluation Metrics.
```

```
from sklearn import metrics
from sklearn.metrics import mean_squared_error
```

```
metrics.r2_score(y,prediction)
```

```
0.5790843515069326
```

```
np.sqrt(mean_squared_error(y,prediction))
```

```
2.077178028336633
```

```
metrics.r2_score(y,prediction)
```

```
0.5790843515069326
```

```
np.sqrt(mean_squared_error(y,prediction))
```

```
2.077178028336633
```

```
print('R_squared:',metrics.r2_score(y,prediction))
print('MSE:',mean_squared_error(y,prediction))
print('RMSE:',np.sqrt(mean_squared_error(y,prediction)))
```

```
R_squared: 0.5790843515069326
```

```
MSE: 4.314668561404462
```

```
RMSE: 2.077178028336633
```

[Colab paid products](#) - [Cancel contracts here](#)

---

✓ 0s completed at 11:10 PM

